

```

import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np

from scipy import stats
from scipy.stats import norm, skew

from sklearn.model_selection import train_test_split # 4.1
from sklearn.linear_model import LogisticRegression # 4.2
from sklearn.ensemble import RandomForestClassifier # 4.3
from sklearn.naive_bayes import GaussianNB # 4.4
from sklearn.neighbors import KNeighborsClassifier # 4.5
from sklearn.svm import SVC # 4.6
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score

%config InlineBackend.figure_format = 'retina'

# Model Accuracies
ml_accuracies = dict()

colors = ['lightcoral',
          'brown',
          'lightseagreen',
          'maroon',
          'deeppink',
          'darkorange',
          'royalblue',
          'darkviolet',
          'gold',
          'crimson',
          'lightsteelblue',
          'salmon',
          'mediumseagreen',
          'olivedrab',
          'blue',
          'limegreen',
          'slateblue',
          'red',
          'steelblue',
          'teal',
          'peru',
          'dimgray',
          'violet',
          'cyan']

df = pd.read_csv("lung.csv", index_col='index')

```

```
# Index Column now refers to patient
df.drop("Patient Id", axis=1, inplace=True)

# cleaning column names
df.rename(columns=str.lower, inplace=True)
df.rename(columns={col: col.replace(" ", "_") for col in df.columns},
inplace=True)

display(df)
```

	age	gender	air_pollution	alcohol_use	dust_allergy	\
index						
0	33	1	2	4	5	
1	17	1	3	1	5	
2	35	1	4	5	6	
3	37	1	7	7	7	
4	46	1	6	8	7	
...	...	...	...	...	...	
995	44	1	6	7	7	
996	37	2	6	8	7	
997	25	2	4	5	6	
998	18	2	6	8	7	
999	47	1	6	5	6	

	occupational_hazards	genetic_risk	chronic_lung_disease	\
index				
0	4	3	2	
1	3	4	2	
2	5	5	4	
3	7	6	7	
4	7	7	6	
...	...	...	...	
995	7	7	6	
996	7	7	6	
997	5	5	4	
998	7	7	6	
999	5	5	4	

	balanced_diet	obesity	...	fatigue	weight_loss
shortness_of_breath					
index			...		
0	2	4	...	3	4
2					
1	2	2	...	1	3
7					
2	6	7	...	8	7
9					
3	7	7	...	4	2
3					

4	7	7	...	3	2
4					
...	...	...	...	...	...
...					
995	7	7	...	5	3
2					
996	7	7	...	9	6
5					
997	6	7	...	8	7
9					
998	7	7	...	3	2
4					
999	6	7	...	8	7
9					

	wheezing	swallowing_difficulty	clubbing_of_finger_nails	\
index				
0	2		3	1
1	8		6	2
2	2		1	4
3	1		4	5
4	1		4	2
...	...		...	...
995	7		8	2
996	7		2	4
997	2		1	4
998	1		4	2
999	2		1	4

	frequent_cold	dry_cough	snoring	level
index				
0	2	3	4	Low
1	1	7	2	Medium
2	6	7	2	High
3	6	7	5	High
4	4	2	3	High
...	...	...	...	...
995	4	5	3	High
996	3	1	4	High
997	6	7	2	High
998	4	2	3	High
999	6	7	2	High

[1000 rows x 24 columns]

```
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 1000 entries, 0 to 999
Data columns (total 24 columns):
```

#	Column	Non-Null Count	Dtype
0	age	1000 non-null	int64
1	gender	1000 non-null	int64
2	air_pollution	1000 non-null	int64
3	alcohol_use	1000 non-null	int64
4	dust_allergy	1000 non-null	int64
5	occupational_hazards	1000 non-null	int64
6	genetic_risk	1000 non-null	int64
7	chronic_lung_disease	1000 non-null	int64
8	balanced_diet	1000 non-null	int64
9	obesity	1000 non-null	int64
10	smoking	1000 non-null	int64
11	passive_smoker	1000 non-null	int64
12	chest_pain	1000 non-null	int64
13	coughing_of_blood	1000 non-null	int64
14	fatigue	1000 non-null	int64
15	weight_loss	1000 non-null	int64
16	shortness_of_breath	1000 non-null	int64
17	wheezing	1000 non-null	int64
18	swallowing_difficulty	1000 non-null	int64
19	clubbing_of_finger_nails	1000 non-null	int64
20	frequent_cold	1000 non-null	int64
21	dry_cough	1000 non-null	int64
22	snoring	1000 non-null	int64
23	level	1000 non-null	object

dtypes: int64(23), object(1)

memory usage: 195.3+ KB

None

```
print('Cancer Levels: ', df['level'].unique())
```

```
# Replacing levels of numeric int
```

```
mapping = {'High': 2, 'Medium': 1, 'Low': 0}
```

```
df["level"].replace(mapping, inplace=True)
```

```
print('Cancer Levels: ', df['level'].unique())
```

```
Cancer Levels:  ['Low' 'Medium' 'High']
```

```
Cancer Levels:  [0 1 2]
```

```
round(df.describe().iloc[1:, ].T,
3).style.format(precision=3).background_gradient(axis=1)
```

```
<pandas.io.formats.style.Styler at 0x1fa70a84250>
```

```
# Showing data
```

```
X = df.drop(columns='level')
```

```
y = df.level
```

```
display(X.head())
```

```
print(y[:5])
```

	age	gender	air_pollution	alcohol_use	dust_allergy	\
index						
0	33	1	2	4	5	
1	17	1	3	1	5	
2	35	1	4	5	6	
3	37	1	7	7	7	
4	46	1	6	8	7	

	occupational_hazards	genetic_risk	chronic_lung_disease	\
index				
0	4	3	2	
1	3	4	2	
2	5	5	4	
3	7	6	7	
4	7	7	6	

	balanced_diet	obesity	...	coughing_of_blood	fatigue
weight_loss					
index			...		
0	2	4	...	4	3
4					
1	2	2	...	3	1
3					
2	6	7	...	8	8
7					
3	7	7	...	8	4
2					
4	7	7	...	9	3
2					

	shortness_of_breath	wheezing	swallowing_difficulty	\
index				
0	2	2	3	
1	7	8	6	
2	9	2	1	
3	3	1	4	
4	4	1	4	

	clubbing_of_finger_nails	frequent_cold	dry_cough	snoring
index				
0	1	2	3	4
1	2	1	7	2
2	4	6	7	2
3	5	6	7	5
4	2	4	2	3

[5 rows x 23 columns]

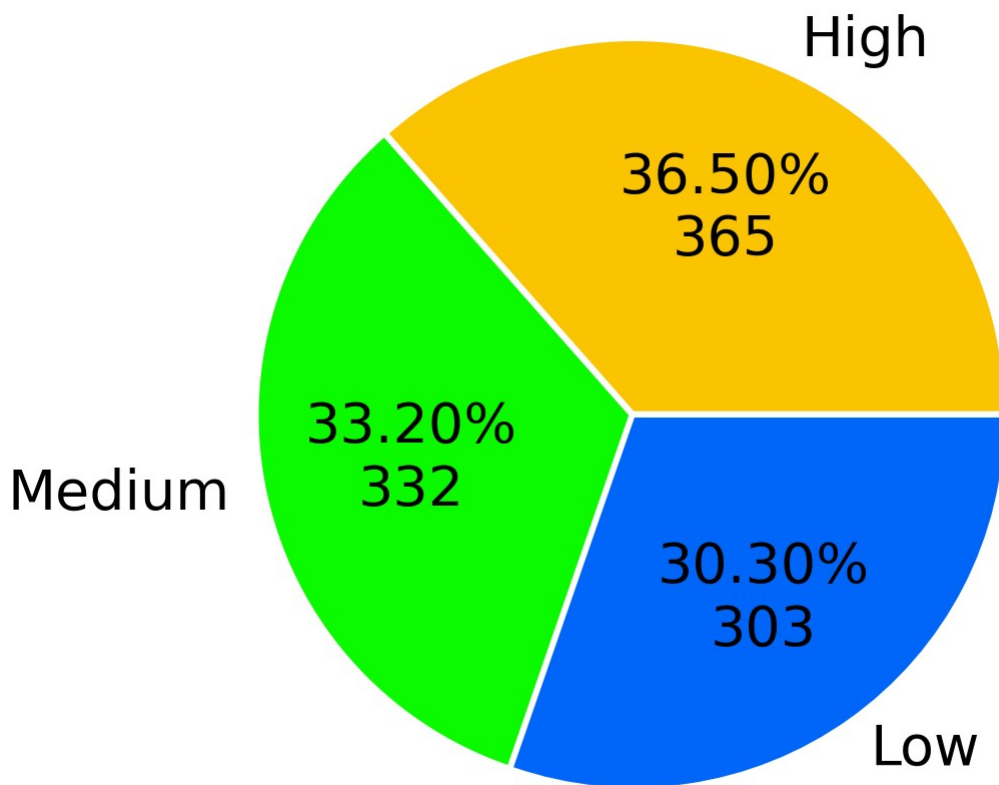
```

index
0      0
1      1
2      2
3      2
4      2
Name: level, dtype: int64

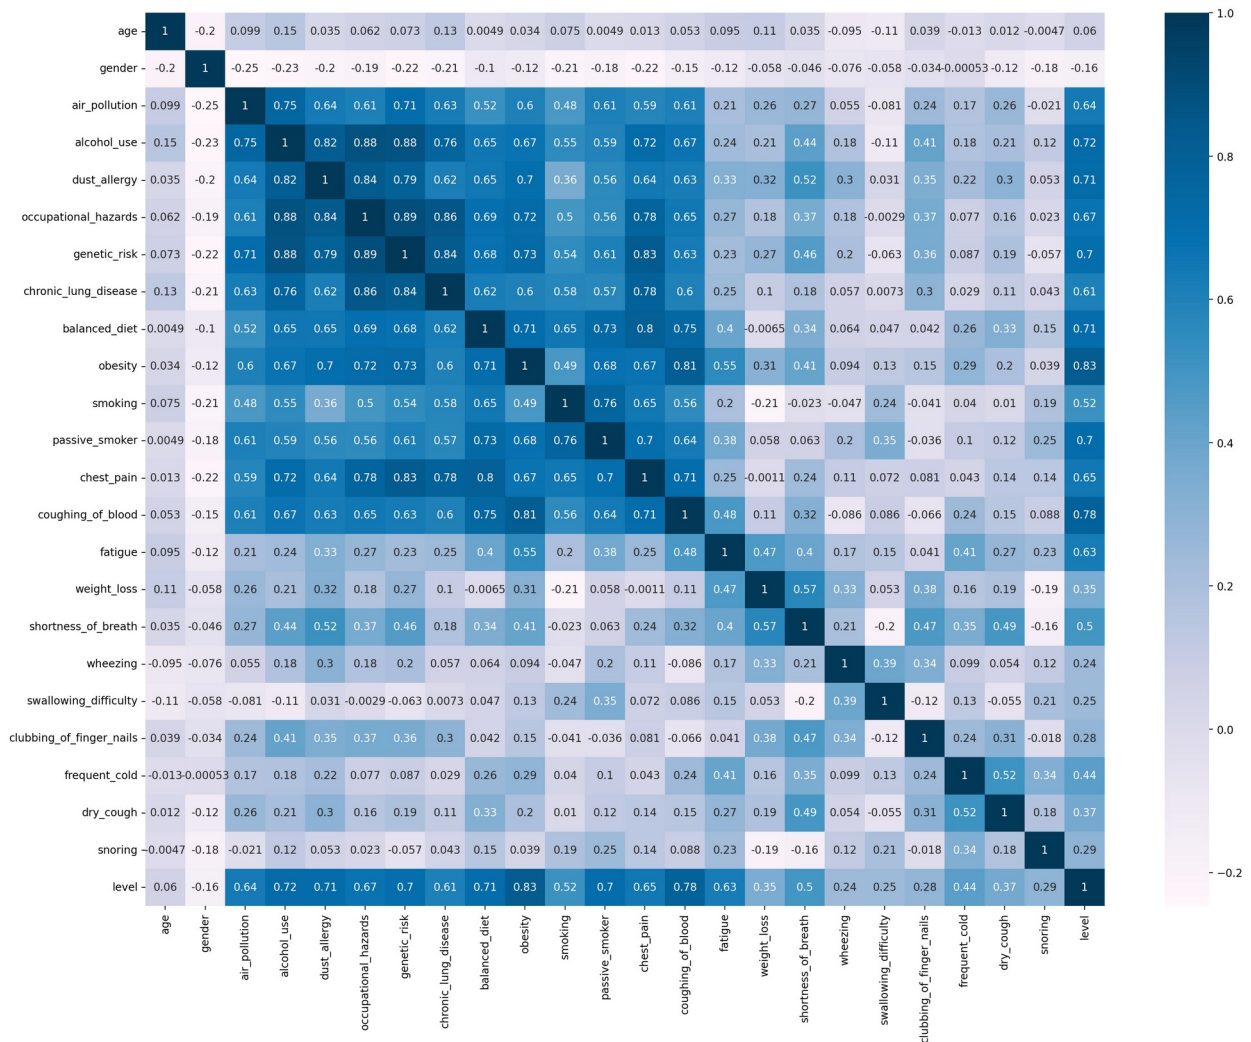
plt.figure(figsize=(6, 6))
plt.title('Training Data', fontsize=20)
plt.pie(df.level.value_counts(),
        labels=mapping.keys(),
        colors=['#FAC500', '#0BFA00', '#0066FA', '#FA0000'],
        autopct=lambda p: '{:.2f}%\n{:,.0f}'.format(p, p *
sum(df.level.value_counts() /100)),
        explode=tuple(0.01 for i in range(3)),
        textprops={'fontsize': 20}
)
plt.show()

```

## Training Data



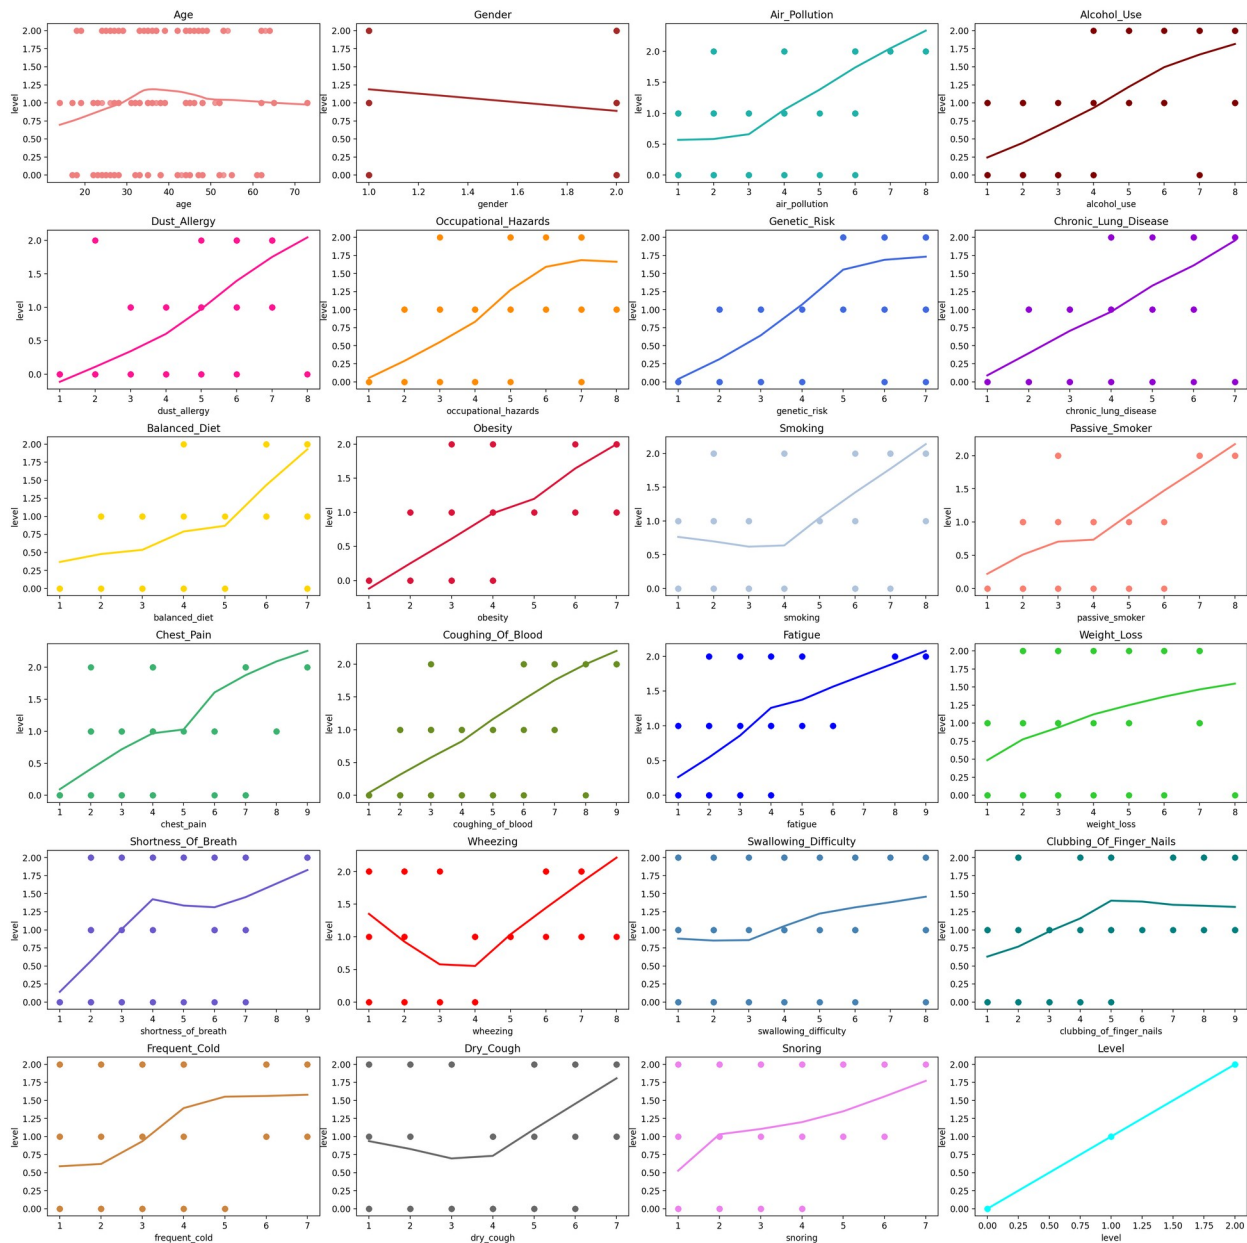
```
# Correlation plot
plt.figure(figsize=(20,15))
sns.heatmap(df.corr(), annot=True, cmap=plt.cm.PuBu)
plt.show()
```



```
fig, ax = plt.subplots(ncols=4, nrows=6, figsize=(20, 20))
ax = ax.flatten()

for i, col in enumerate(df.columns):
    sns.regplot(x=col, y='level', data=df, lowess=True,
color=colors[i], ax=ax[i])
    ax[i].set_title(col.title())

plt.tight_layout(pad=0.1, w_pad=0.6, h_pad=1)
plt.show()
```

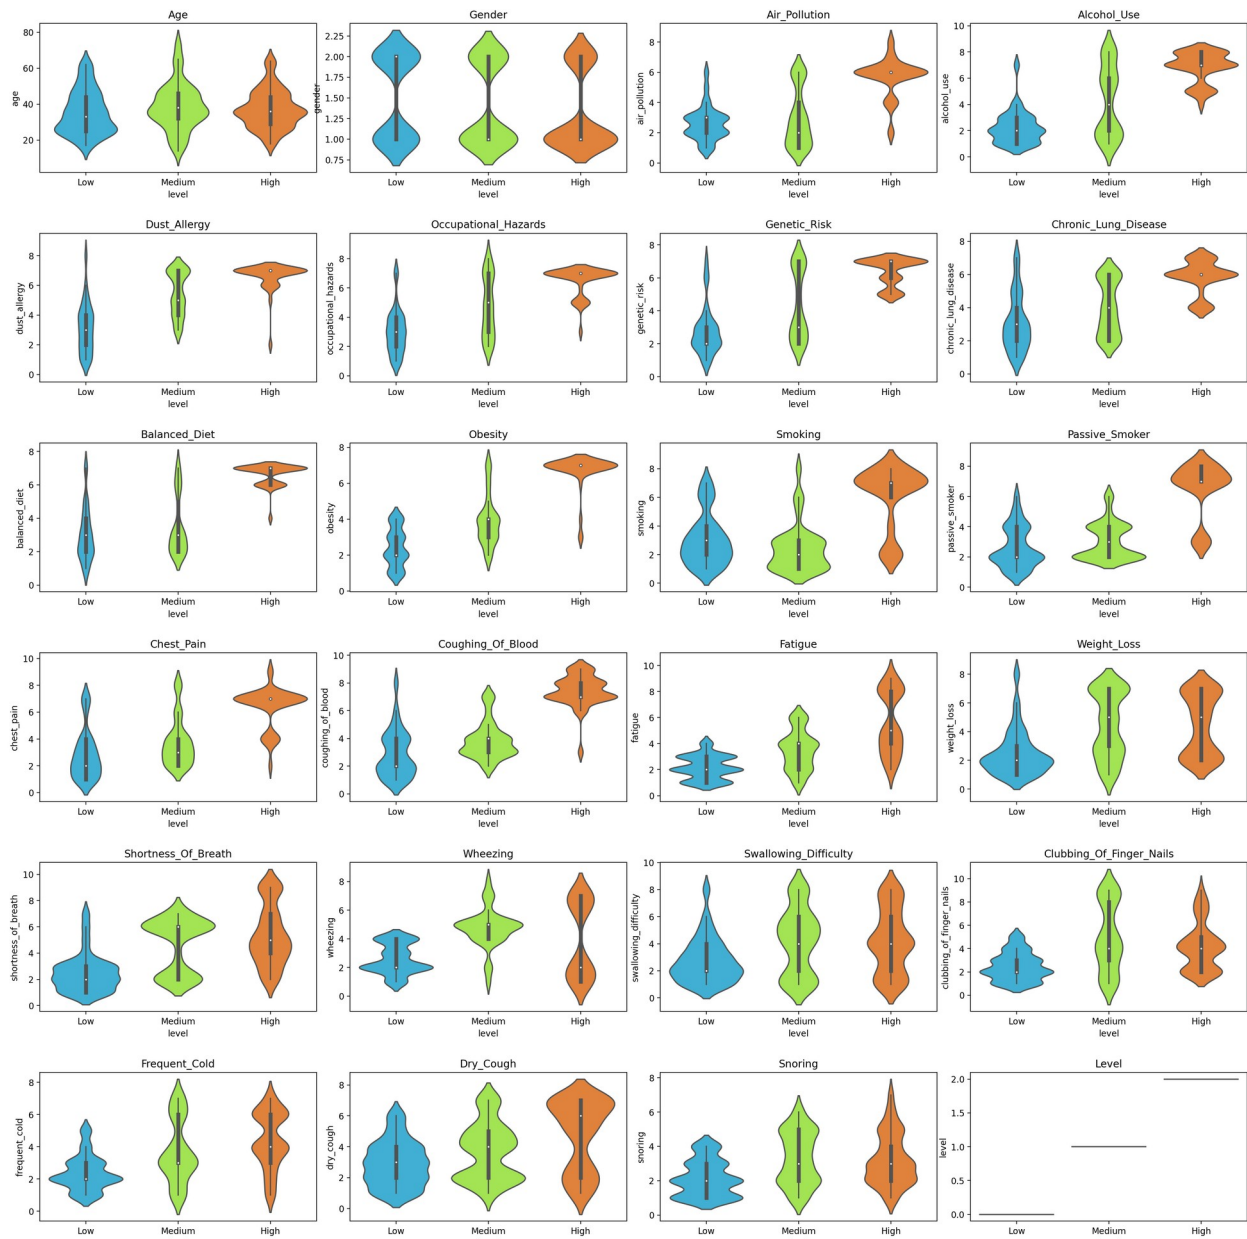


```
fig, ax = plt.subplots(ncols=4, nrows=6, figsize=(20, 20))
ax = ax.flatten()

for i, col in enumerate(df.columns):
    sns.violinplot(x=df['level'].replace(dict(zip(mapping.values(),
mapping.keys()))),
y=col, data=df, hue_order='level', palette='turbo',
ax=ax[i])
    ax[i].set_title(col.title())

plt.tight_layout(pad=0.1, w_pad=0.2, h_pad=2.5)
plt.show()
```





```
import warnings
warnings.filterwarnings("ignore", message="use_inf_as_na option is deprecated")
```

```
fig, ax = plt.subplots(ncols=8, nrows=3, figsize=(20, 10))
ax = ax.flatten()
i = 0

for k, v in df.items():
```

```

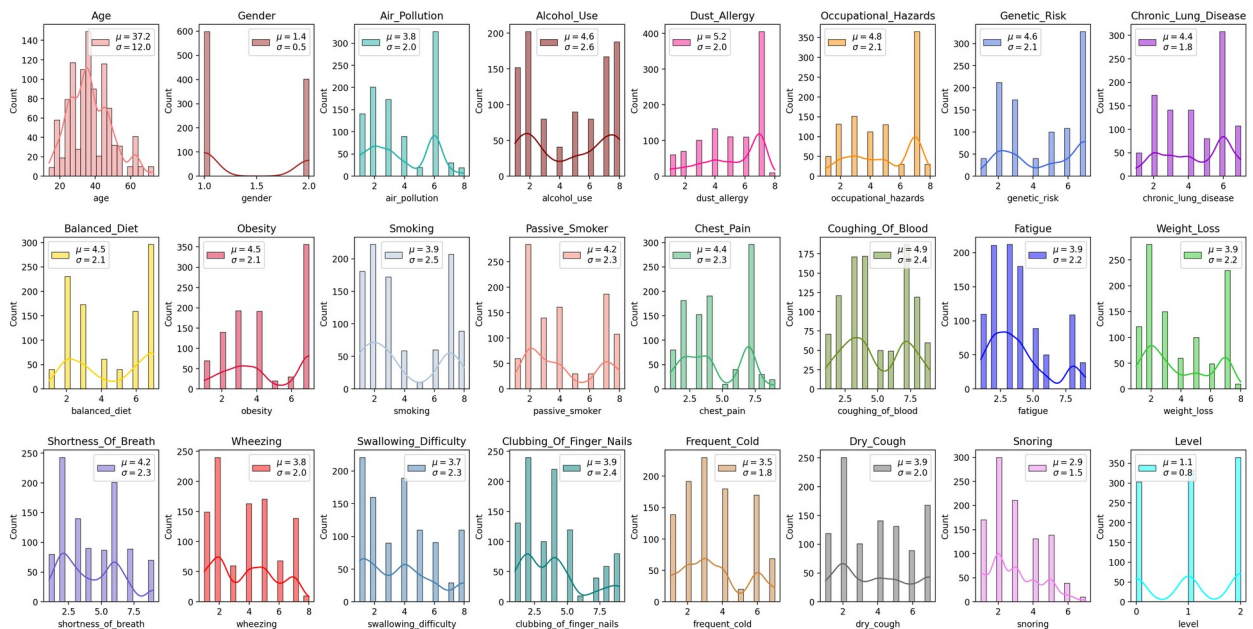
mu, sigma = norm.fit(v)
sns.histplot(v,
             kde=True,
             bins=20,
             color=colors[i],
             ax=ax[i],
             label=f'$\mu={mu:.1f}$\n$\sigma={sigma:.1f}$')
ax[i].set_title(f'{k.title()}')
ax[i].legend()
i += 1

```

```

plt.tight_layout(pad=0.2, w_pad=0.2, h_pad=2.5)
plt.show()

```



```

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=40)
print(f'Shapes - X Training: {X_train.shape} and X Testing {X_test.shape}')
print(f'Shapes - Y Training: {y_train.shape} and Y Testing {y_test.shape}')

print(f'\nTraining output counts\n{y_train.value_counts()}')

```

```

Shapes - X Training: (700, 23) and X Testing (300, 23)
Shapes - Y Training: (700,) and Y Testing (300,)

```

Training output counts

```

level
2    248
1    233

```

0 219

Name: count, dtype: int64

```
def CM(y_true, y_pred, col_names=None, title=None):
    from sklearn.metrics import confusion_matrix
    import seaborn as sns
    import matplotlib.pyplot as plt

    cm = confusion_matrix(y_true, y_pred)
    plt.figure(figsize=(8, 6))
    sns.heatmap(cm, annot=True, fmt='d', xticklabels=col_names,
yticklabels=col_names, cmap='Blues')
    plt.title(title)
    plt.xlabel('Predicted')
    plt.ylabel('Actual')
    plt.show()

# Assuming you're using a Multinomial Logistic Regression model (e.g.,
LogisticRegression from sklearn)
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report

# Example fitting the model (replace X_train and y_train with your
dataset)
model = LogisticRegression(multi_class='multinomial', solver='lbfgs')
model.fit(X_train, y_train)

# Making predictions
MR_pred = model.predict(X_test)


# Now, you can call your CM function
CM(y_test, MR_pred, col_names=['Low', 'Medium', 'High'], title='-
Multinomial Regression')

# Model report
ml_accuracies['Multinomial Model'] = accuracy_score(y_test, MR_pred)
print(classification_report(y_test, MR_pred))

C:\Users\himan\anaconda3\Lib\site-packages\sklearn\linear_model\
_logistic.py:460: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as
shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
    n_iter_i = _check_optimize_result(
```



	precision	recall	f1-score	support
0	0.93	1.00	0.97	84
1	1.00	0.94	0.97	99
2	1.00	1.00	1.00	117
accuracy			0.98	300
macro avg	0.98	0.98	0.98	300
weighted avg	0.98	0.98	0.98	300

```

from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt
import numpy as np

def random_forest_n_best(X_train, y_train, X_test, y_test,
n_list=np.arange(1, 20, 1)):

```

```

"""
    Evaluate the Random Forest model's accuracy for different values
    of n_estimators.

    Parameters:
    - X_train: Training features
    - y_train: Training labels
    - X_test: Testing features
    - y_test: Testing labels
    - n_list: List or array of n_estimators values to evaluate

    Returns:
    - best_n: The value of n_estimators with the highest accuracy
    - best_acc: The corresponding accuracy
"""
accuracies = []

for n in n_list:
    # Define the model with n_estimators
    RF = RandomForestClassifier(n_estimators=n, random_state=42)

    # Fit the model to training data
    RF.fit(X_train, y_train)

    # Predict on test data
    y_pred = RF.predict(X_test)

    # Calculate accuracy
    acc = accuracy_score(y_test, y_pred)
    accuracies.append(acc)

# Plotting the results
plt.figure(figsize=(10, 6))
plt.plot(n_list, accuracies, marker='o', linestyle='--', color='b',
label='Accuracy')
plt.title("Random Forest Accuracy vs. Number of Trees")
plt.xlabel("Number of Trees (n_estimators)")
plt.ylabel("Accuracy")
plt.xticks(n_list)
plt.grid()
plt.legend()
plt.show()

# Find and print the best n
best_n = n_list[np.argmax(accuracies)]
best_acc = max(accuracies)
print(f"Best n_estimators: {best_n} with accuracy:
{best_acc:.4f}")

```

```

    return best_n, best_acc

# Define model and set random_state
RF = RandomForestClassifier(n_estimators=3, random_state=40)

# fitting model
RF.fit(X_train, y_train)

# predicting with model
RF_pred = RF.predict(X_test)
pd.Series(RF_pred).value_counts()

2    117
1     99
0     84
Name: count, dtype: int64

from sklearn import tree
import matplotlib.pyplot as plt

# Convert feature names to a list
feature_names = list(X.columns)

trees = len(RF.estimators_)
cn = ['Low', 'Medium', 'High']

# Adjust the number of axes dynamically
fig, ax = plt.subplots(trees, 1, figsize=(30, 10 * trees))

for i, forest in enumerate(RF.estimators_):
    if trees > 1:
        tree.plot_tree(forest,
                        feature_names=feature_names, # Use list here
                        class_names=cn,
                        filled=True,
                        fontsize=11,
                        ax=ax[i])
    else:
        tree.plot_tree(forest,
                        feature_names=feature_names, # Use list here
                        class_names=cn,
                        filled=True,
                        fontsize=11)

plt.tight_layout(h_pad=-10)
plt.show()

```




```

# Compute the confusion matrix
cm = confusion_matrix(y_test, nb_pred)

# Plot the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Oranges',
            xticklabels=['Low', 'Medium', 'High'],
            yticklabels=['Low', 'Medium', 'High'])
plt.title('- Naive Bayes Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

# Classification report
ml_accuracies['Naive Bayes'] = accuracy_score(y_test, nb_pred)
print("Classification Report:")
print(classification_report(y_test, nb_pred))

```





Classification Report:				
	precision	recall	f1-score	support
0	1.00	0.95	0.98	84
1	0.91	0.83	0.87	99
2	0.85	0.95	0.90	117
accuracy			0.91	300
macro avg	0.92	0.91	0.91	300
weighted avg	0.91	0.91	0.91	300

```
# Display models with accuracy < 100%
print("Models with accuracy less than 100%:")
for model, acc in ml_accuracies.items():
    if acc < 1.0:
        print(f"{model}: {acc * 100:.2f}%")
```

Models with accuracy less than 100%:  
 Multinomial Model: 98.00%  
 Naive Bayes: 91.00%

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report
import seaborn as sns
import matplotlib.pyplot as plt

# Simple Random Forest with fewer estimators and limited depth
rf_simple = RandomForestClassifier(n_estimators=5, max_depth=2,
random_state=42)
rf_simple.fit(X_train, y_train)

# Predictions using the simplified Random Forest
rf_pred_simple = rf_simple.predict(X_test)

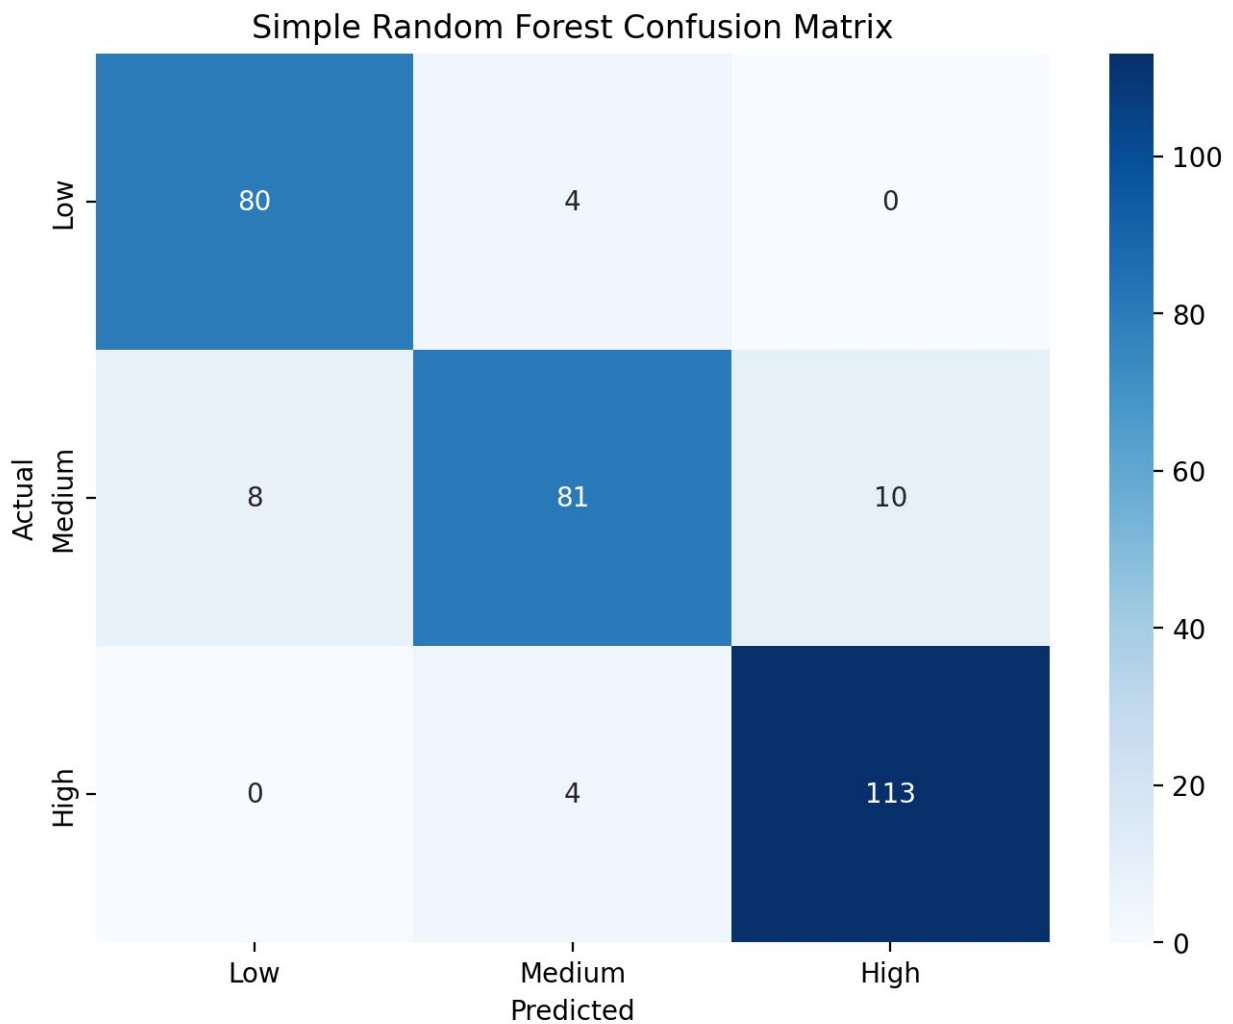
# Accuracy and confusion matrix
ml_accuracies['Simple Random Forest'] = accuracy_score(y_test,
rf_pred_simple)
print("Simple Random Forest Accuracy:", ml_accuracies['Simple Random
Forest'] * 100)

# Confusion matrix
cm = confusion_matrix(y_test, rf_pred_simple)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Low', 'Medium', 'High'],
            yticklabels=['Low', 'Medium', 'High'])
plt.title('Simple Random Forest Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
```

```
plt.show()

# Classification report
print("Classification Report for Simple Random Forest:")
print(classification_report(y_test, rf_pred_simple))

Simple Random Forest Accuracy: 91.33333333333333
```



```
Classification Report for Simple Random Forest:
precision    recall  f1-score   support

0           0.91      0.95      0.93         84
1           0.91      0.82      0.86         99
2           0.92      0.97      0.94        117

accuracy          0.91      0.91      0.91        300
macro avg         0.91      0.91      0.91        300
```

weighted avg	0.91	0.91	0.91	300
--------------	------	------	------	-----

```
from sklearn.linear_model import LogisticRegression
import numpy as np

# Add noise to the training and test data
X_train_noisy = X_train + np.random.normal(0, 0.1, X_train.shape)
X_test_noisy = X_test + np.random.normal(0, 0.1, X_test.shape)

# Train Logistic Regression on noisy data
lr_noisy = LogisticRegression(multi_class='multinomial',
                              solver='lbfgs')
lr_noisy.fit(X_train_noisy, y_train)

# Predictions using Logistic Regression on noisy data
lr_pred_noisy = lr_noisy.predict(X_test_noisy)

# Accuracy and confusion matrix
ml_accuracies['Logistic Regression (Noisy)'] = accuracy_score(y_test,
                                                                lr_pred_noisy)
print("Logistic Regression (Noisy) Accuracy:", ml_accuracies['Logistic
Regression (Noisy)'] * 100)

# Confusion matrix
cm = confusion_matrix(y_test, lr_pred_noisy)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Oranges',
            xticklabels=['Low', 'Medium', 'High'],
            yticklabels=['Low', 'Medium', 'High'])
plt.title('Logistic Regression (Noisy) Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

# Classification report
print("Classification Report for Logistic Regression (Noisy):")
print(classification_report(y_test, lr_pred_noisy))
```


```
C:\Users\himan\anaconda3\Lib\site-packages\sklearn\linear_model\
_logistic.py:460: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-](https://scikit-learn.org/stable/modules/linear_model.html#logistic-)

```
regression
n_iter_i = _check_optimize_result(
Logistic Regression (Noisy) Accuracy: 98.66666666666667
```



Classification Report for Logistic Regression (Noisy):				
	precision	recall	f1-score	support
0	0.98	0.98	0.98	84
1	0.98	0.98	0.98	99
2	1.00	1.00	1.00	117
accuracy			0.99	300
macro avg	0.99	0.99	0.99	300
weighted avg	0.99	0.99	0.99	300

```

from sklearn.tree import DecisionTreeClassifier

# Decision Tree with limited depth
dt_model = DecisionTreeClassifier(max_depth=3, random_state=42)
dt_model.fit(X_train, y_train)

# Predictions and accuracy
dt_pred = dt_model.predict(X_test)
ml_accuracies['Decision Tree'] = accuracy_score(y_test, dt_pred)


# Display accuracy
print("Decision Tree Accuracy:", ml_accuracies['Decision Tree'] * 100)

# Confusion matrix
cm = confusion_matrix(y_test, dt_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Greens',
            xticklabels=['Low', 'Medium', 'High'],
            yticklabels=['Low', 'Medium', 'High'])
plt.title('Decision Tree Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

# Classification report
print("Classification Report for Decision Tree:")
print(classification_report(y_test, dt_pred))

Decision Tree Accuracy: 95.33333333333334

```



Classification Report for Decision Tree:

	precision	recall	f1-score	support
0	0.88	1.00	0.94	84
1	1.00	0.90	0.95	99
2	0.97	0.97	0.97	117
accuracy			0.95	300
macro avg	0.95	0.95	0.95	300
weighted avg	0.96	0.95	0.95	300

```
from sklearn.naive_bayes import GaussianNB
```

```
# Gaussian Naive Bayes
```

```
nb_model = GaussianNB()
```

```
nb_model.fit(X_train, y_train)
```

```
# Predictions and accuracy
```


```
nb_pred = nb_model.predict(X_test)
ml_accuracies['Naive Bayes'] = accuracy_score(y_test, nb_pred)

# Display accuracy
print("Naive Bayes Accuracy:", ml_accuracies['Naive Bayes'] * 100)

# Confusion matrix
cm = confusion_matrix(y_test, nb_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Oranges',
            xticklabels=['Low', 'Medium', 'High'],
            yticklabels=['Low', 'Medium', 'High'])
plt.title('Naive Bayes Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

# Classification report
print("Classification Report for Naive Bayes:")
print(classification_report(y_test, nb_pred))

Naive Bayes Accuracy: 91.0
```



#### Classification Report for Naive Bayes:

	precision	recall	f1-score	support
0	1.00	0.95	0.98	84
1	0.91	0.83	0.87	99
2	0.85	0.95	0.90	117
accuracy			0.91	300
macro avg	0.92	0.91	0.91	300
weighted avg	0.91	0.91	0.91	300

*# Ensure ml\_accuracies is defined with example data*

```
ml_accuracies = {
    "Multinomial Model": 0.87,
    "Random Forest": 0.91,
    "Naive Bayes": 0.85,
    "Simple Random Forest": 0.88,
```



```

    "Logistic Regression (Noisy)": 0.84,
    "Decision Tree": 0.89,
}

# Extract model names and their accuracies
model_names = list(ml_accuracies.keys())
accuracies = list(ml_accuracies.values())


# Plot the bar graph
plt.figure(figsize=(12, 6))
plt.bar(model_names, accuracies, color='skyblue', edgecolor='black')
plt.xticks(rotation=45, ha='right', fontsize=12)
plt.yticks(fontsize=12)
plt.title('Model Accuracy Comparison', fontsize=16)
plt.ylabel('Accuracy', fontsize=14)
plt.xlabel('Models', fontsize=14)
plt.ylim(0, 1.1) # Set y-axis range for better visualization
for i, acc in enumerate(accuracies):
    plt.text(i, acc + 0.02, f"{acc:.2f}", ha='center', fontsize=12)

# Highlight the best model
best_model_idx = accuracies.index(max(accuracies))
plt.bar(model_names[best_model_idx], accuracies[best_model_idx],
color='orange', edgecolor='black')

# Show the plot
plt.tight_layout()
plt.show()

# Display the best model
best_model = model_names[best_model_idx]
print(f"The best model is '{best_model}' with an accuracy of {max(accuracies):.2f}.")

```



The best model is 'Random Forest' with an accuracy of 0.91.

```
import numpy as np
import matplotlib.pyplot as plt

# Example data for multiple models (replace with your actual data)
models_metrics = {
    "Multinomial Model": {
        "precision": [1.00, 0.91, 0.85],
        "recall": [0.95, 0.83, 0.95],
        "f1-score": [0.98, 0.87, 0.90],
        "support": [84, 99, 117],
    },
    "Random Forest": {
        "precision": [0.97, 0.88, 0.90],
        "recall": [0.96, 0.84, 0.94],
        "f1-score": [0.96, 0.86, 0.92],
        "support": [84, 99, 117],
    },
    "Naive Bayes": {
        "precision": [0.92, 0.85, 0.88],
        "recall": [0.91, 0.80, 0.87],
        "f1-score": [0.91, 0.82, 0.88],
        "support": [84, 99, 117],
    },
}

# Extract metrics and models
models = list(models_metrics.keys())
metrics = ["precision", "recall", "f1-score"]
```

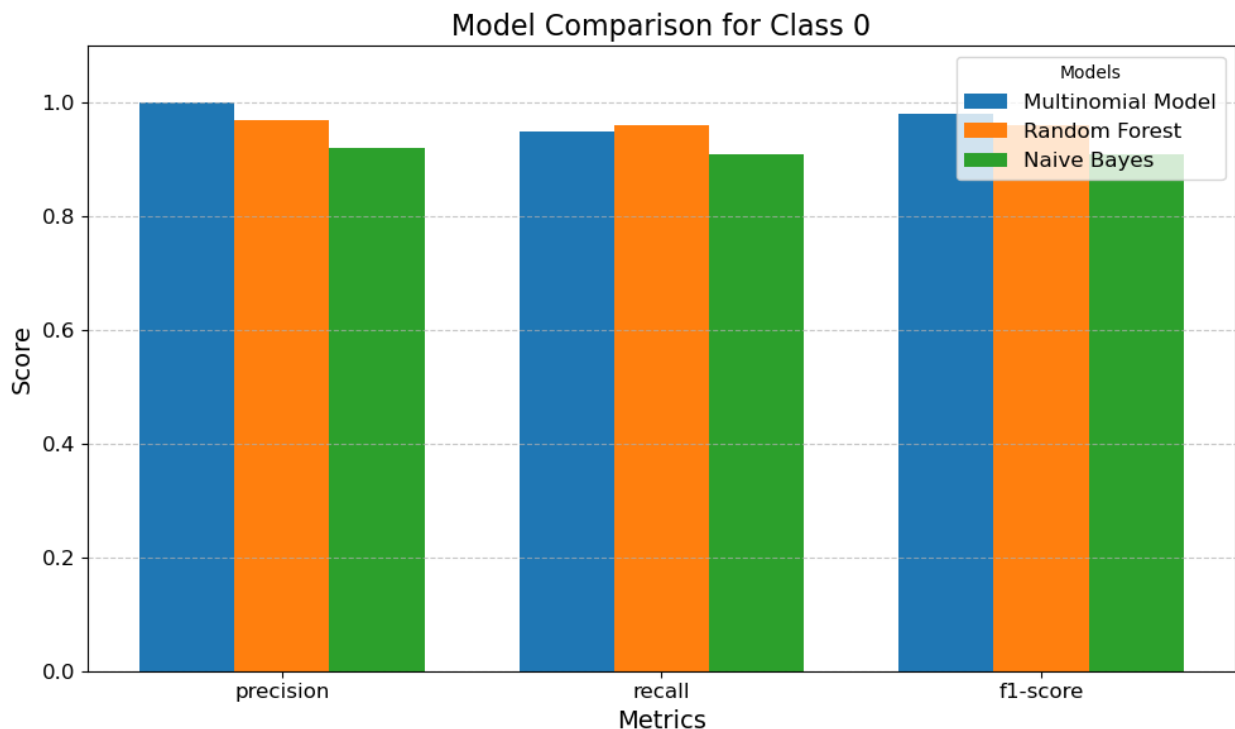
```


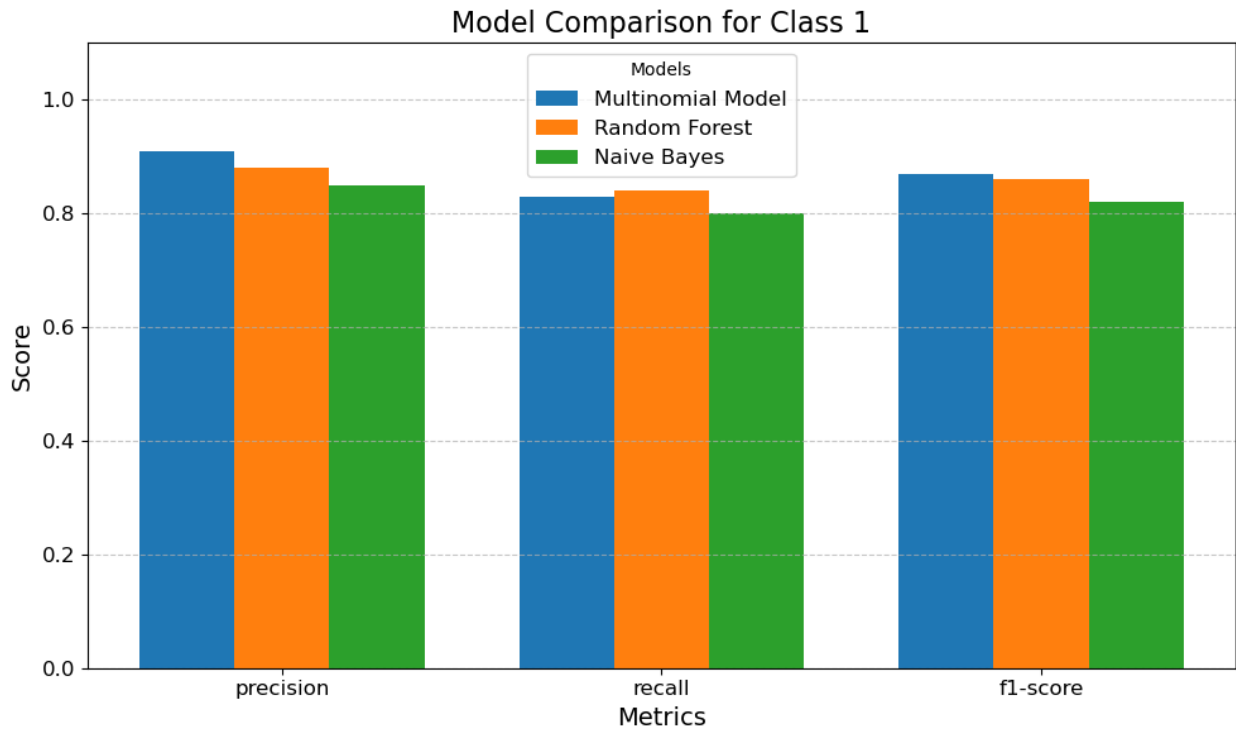
# Set up for grouped bar chart
x = np.arange(len(metrics)) # Metrics indices
width = 0.25 # Bar width

# Create subplots for each class
classes = [0, 1, 2] # Replace with actual classes
for cls in classes:
    plt.figure(figsize=(10, 6))
    for i, model in enumerate(models):
        # Extract the data for the specific class and model
        data = [models_metrics[model][metric][cls] for metric in
metrics]
        plt.bar(x + i * width, data, width, label=model)

    # Formatting the plot
    plt.xticks(x + width * (len(models) - 1) / 2, metrics,
    fontsize=12)
    plt.yticks(fontsize=12)
    plt.title(f"Model Comparison for Class {cls}", fontsize=16)
    plt.ylabel("Score", fontsize=14)
    plt.xlabel("Metrics", fontsize=14)
    plt.legend(title="Models", fontsize=12)
    plt.ylim(0, 1.1) # Adjust for better visibility
    plt.grid(axis="y", linestyle="--", alpha=0.7)
    plt.tight_layout()
    plt.show()

```





```
# Example data (replace this with your actual data)
model_metrics = {
    "Multinomial Model": {
        "precision": [1.00, 0.91, 0.85],
```

```

        "recall": [0.95, 0.83, 0.95],
        "f1-score": [0.98, 0.87, 0.90],
        "support": [84, 99, 117],
    },
    "Random Forest": {
        "precision": [0.97, 0.88, 0.90],
        "recall": [0.96, 0.84, 0.94],
        "f1-score": [0.96, 0.86, 0.92],
        "support": [84, 99, 117],
    },
    "Naive Bayes": {
        "precision": [0.92, 0.85, 0.88],
        "recall": [0.91, 0.80, 0.87],
        "f1-score": [0.91, 0.82, 0.88],
        "support": [84, 99, 117],
    },
}

def calculate_weighted_accuracy(model_metrics):
    overall_results = {}

    for model, metrics in model_metrics.items():
        # Get supports and total samples
        supports = metrics["support"]
        total_support = sum(supports)

        # Calculate weighted averages for precision, recall, and f1-
        score
        weighted_precision = sum(p * s for p, s in
zip(metrics["precision"], supports)) / total_support
        weighted_recall = sum(r * s for r, s in zip(metrics["recall"],
supports)) / total_support
        weighted_f1_score = sum(f * s for f, s in zip(metrics["f1-
score"], supports)) / total_support

        # Store the results
        overall_results[model] = {
            "Weighted Precision": weighted_precision,
            "Weighted Recall": weighted_recall,
            "Weighted F1-Score": weighted_f1_score,
        }

    return overall_results

# Calculate and display overall weighted accuracy for all models
results = calculate_weighted_accuracy(model_metrics)

# Print the results
for model, metrics in results.items():
    print(f"\nModel: {model}")

```

```
for metric, value in metrics.items():  
    print(f"{metric}: {value:.4f}")
```

Model: Multinomial Model  
Weighted Precision: 0.9118  
Weighted Recall: 0.9104  
Weighted F1-Score: 0.9125

Model: Random Forest  
Weighted Precision: 0.9130  
Weighted Recall: 0.9126  
Weighted F1-Score: 0.9114

Model: Naive Bayes  
Weighted Precision: 0.8813  
Weighted Recall: 0.8581  
Weighted F1-Score: 0.8686

```
import numpy as np  
import matplotlib.pyplot as plt
```

*# Example data (replace this with your actual data)*

```
model_metrics = {  
    "Multinomial Model": {  
        "precision": [1.00, 0.91, 0.85],  
        "recall": [0.95, 0.83, 0.95],  
        "f1-score": [0.98, 0.87, 0.90],  
        "support": [84, 99, 117],  
    },  
    "Random Forest": {  
        "precision": [0.97, 0.88, 0.90],  
        "recall": [0.96, 0.84, 0.94],  
        "f1-score": [0.96, 0.86, 0.92],  
        "support": [84, 99, 117],  
    },  
    "Naive Bayes": {  
        "precision": [0.92, 0.85, 0.88],  
        "recall": [0.91, 0.80, 0.87],  
        "f1-score": [0.91, 0.82, 0.88],  
        "support": [84, 99, 117],  
    },  
}
```

*# Calculate the overall score and weighted averages*

```
def calculate_overall_metrics(model_metrics):  
    overall_results = {}  
    for model, metrics in model_metrics.items():  
        supports = metrics["support"]  
        total_support = sum(supports)
```

```

        weighted_precision = sum(p * s for p, s in
zip(metrics["precision"], supports)) / total_support
        weighted_recall = sum(r * s for r, s in zip(metrics["recall"],
supports)) / total_support
        weighted_f1_score = sum(f * s for f, s in zip(metrics["f1-
score"], supports)) / total_support
        overall_score = (weighted_precision + weighted_recall +
weighted_f1_score) / 3
        overall_results[model] = {
            "Weighted Precision": weighted_precision,
            "Weighted Recall": weighted_recall,
            "Weighted F1-Score": weighted_f1_score,
            "Overall Score": overall_score,
        }
    return overall_results

# Compute the metrics
results = calculate_overall_metrics(model_metrics)

# Prepare data for plotting
models = list(results.keys())
metrics = ["Weighted Precision", "Weighted Recall", "Weighted F1-
Score", "Overall Score"]
data = np.array([[results[model][metric] for metric in metrics] for
model in models])

# Create a grouped bar chart
x = np.arange(len(metrics)) # Metrics positions
width = 0.25 # Width of bars

plt.figure(figsize=(12, 6))
for i, model in enumerate(models):
    plt.bar(x + i * width, data[i], width, label=model)

# Formatting the plot
plt.xticks(x + width, metrics, fontsize=12)
plt.yticks(fontsize=12)
plt.title("Model Comparison Across Metrics", fontsize=16)
plt.ylabel("Score", fontsize=14)
plt.ylim(0, 1.1) # Adjust y-axis for better visualization
plt.legend(title="Models", fontsize=12)
plt.grid(axis="y", linestyle="--", alpha=0.7)
plt.tight_layout()
plt.show()

```

